



COMPUTER SCIENCE AND DATA ANALYTICS
Course: CSCI 6444 Intro to Big Data Analytics

Class project #1

R and Graph Analytics

Student(s): Group 6
Anar Shikhaliyev
Eljan Mammadli

Instructor: Dr. Abzatdin Adamov

Baku 2023

1. Data Set

In this assignment, we have a network of all the incoming and outgoing email of the research institution communication between European countries. We have 3,038,531 emails in total, distributed among 287,755 different email accounts. We only have a complete email graph for 1,258 of the research institution's email addresses. Also, within the dataset's time frame, 34,203 email addresses sent and received email. All other email addresses are either invalid, typographically incorrect, or spam.

2. Install the igraph package from one of the CRAN mirrors

In order to install the igraph, we need to follow these steps:

1. Open RStudio
2. Type following command in the R console: `install.packages("igraph")`

```
> install.packages("igraph")
also installing the dependencies 'magrittr', 'pkgconfig', 'rlang'

trying URL 'https://cran.rstudio.com/bin/macosx/big-sur-arm64/contrib/4.2/magrittr_2.0.3.tgz'
Content type 'application/x-gzip' length 231251 bytes (225 KB)
=====
downloaded 225 KB

trying URL 'https://cran.rstudio.com/bin/macosx/big-sur-arm64/contrib/4.2/pkgconfig_2.0.3.tgz'
Content type 'application/x-gzip' length 17697 bytes (17 KB)
=====
downloaded 17 KB

trying URL 'https://cran.rstudio.com/bin/macosx/big-sur-arm64/contrib/4.2/rlang_1.1.0.tgz'
Content type 'application/x-gzip' length 1867245 bytes (1.8 MB)
=====
downloaded 1.8 MB

trying URL 'https://cran.rstudio.com/bin/macosx/big-sur-arm64/contrib/4.2/igraph_1.4.1.tgz'
Content type 'application/x-gzip' length 8231069 bytes (7.8 MB)
=====
downloaded 7.8 MB

The downloaded binary packages are in
  /var/folders/wg/crkb368569v_9g9rddhss1jc0000gn/T//Rtmp9l8QEq/downloaded_packages
>
```

Picture 1. Installing igraph

3. R did not prompt any CRAN mirror selection because, it installed packages from default CRAN mirror in R configuration which we can access it by:

```
> getOption("repos")
CRAN
"https://cran.rstudio.com/"
attr(,"RStudio")
[1] TRUE
> |
```

3. Experiment with some of the functions that shown in the Introduction to Graph

Analytics document on Blackboard on the graph generated from the data set. Present the results in your write-up.

We have experimented this R functions on the given dataset:

- Vcount(), V() and ecount(), E() - These functions help to gain information about the vectors and edges of the graph respectively.

```
> vcount(email_graph)
[1] 265214
> V(email_graph)
+ 265214/265214 vertices, named, from bece4a0:
[1] 0 1 2 3 5 6 7 8 9 10 11 12 14 15 16 18 19 20 21 22 23 24 25 26 27 28 29
[28] 30 31 32 33 34 35 36 37 39 40 41 42 43 44 45 46 47 48 49 50 51 53 54 55 56 57 58
[55] 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 76 77 78 79 80 81 82 83 84 86 87
[82] 88 90 92 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108 109 111 113 114 115 116 118 119 120
[109] 121 123 124 125 126 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144 145 146 147 148
[136] 149 150 151 152 153 155 156 158 159 160 162 163 164 165 166 167 168 170 171 172 174 175 177 178 179 181 182
[163] 183 184 185 186 187 189 191 192 193 194 195 196 197 199 200 201 202 203 204 205 206 207 208 209 210 211 212
[190] 214 215 216 217 218 219 220 221 222 223 225 226 228 229 231 232 233 234 236 237 238 239 240 241 242 243 244
[217] 246 247 248 249 250 251 252 253 254 255 256 259 260 261 263 264 265 266 268 269 270 271 272 273 274 275 276
[244] 278 279 280 281 283 284 285 286 287 288 289 291 292 293 294 295 296 297 298 299 300 301 302 303 304 305 306
+ ... omitted several vertices
```

```

> ecount(email_graph)
[1] 420045
> E(email_graph)
+ 420045/420045 edges from bece4a0 (vertex names):
[1] 0->1    0->4    0->5    0->8    0->11   0->20   0->48   0->130  0->160  0->430  0->668  0->736  0->3612 0->4252 0->16687
[16] 1->1    1->44   1->50   1->56   1->98   1->99   1->106  1->146  1->147  1->149  1->158  1->171  1->175  1->184  1->206
[31] 1->259  1->333  1->336  1->392  1->397  1->406  1->422  1->446  1->457  1->585  1->602  1->620  1->640  1->732  1->733
[46] 1->779  1->841  1->1033 1->1118 1->1261 1->1262 1->1290 1->1370 1->1425 1->1458 1->1515 1->1518 1->1521 1->1546 1->1619
[61] 1->1623 1->1776 1->1803 1->1966 1->1969 1->2014 1->2037 1->2058 1->2244 1->2356 1->2558 1->2874 1->2924 1->3449 1->4200
[76] 1->4681 1->5357 1->5592 1->5726 1->6044 1->6119 1->6962 1->7719 1->7723 1->7817 1->7998 1->8027 1->8109 1->8244 1->9394
[91] 1->14402 1->14507 1->15162 1->15439 1->16393 1->17166 1->19459 1->19578 1->20587 1->20833 1->21290 1->21442 1->22604 1->23384 1->23659
[106] 1->23783 1->23927 1->24137 1->24339 1->24466 1->26034 1->26354 1->27008 1->29534 1->30331 1->33287 1->33976 1->34212 1->34609 1->34630
[121] 1->34638 1->34710 1->34792 1->38060 1->38300 1->39437 1->39527 1->39532 1->39811 1->40430 1->40996 1->41040 1->42064 1->44877 1->45382
[136] 1->45461 1->48364 1->55144 1->55254 1->55820 1->59020 1->60576 1->63349 1->68370 1->69668 1->71303 1->74996 1->75557 1->75912 1->76060
+ ... omitted several edges

```

- Density - proportion between the number of edges and the number of potential edges is known as a graph's density.

```

> edge_density(email_graph)
[1] 5.971791e-06

```

- Degree - shows the degree of every node inside the graph. Since our data is too big we will not show the actual output.

Since our data is too large some of the commands were taking too much time and resource to execute. For that reason, it might be helpful to simplify the graph when working with big networks like ours to minimize the number of nodes and edges, making it easier to manipulate. So, we need to simplify the graph in order to execute those commands.

- Simplify and is_simple - eliminates the loop as well as several edges from a graph. As we can see from the output below, there is a slight decrease in number of edges after removing loops and it is verified by the is_simple function.

```

> simple_email = simplify(email_graph)
> ecount(email_graph)
[1] 420045
> ecount(simple_email)
[1] 418956
> is_simple(email_graph)
[1] FALSE
> is_simple(simple_email)
[1] TRUE

```

But doing only this was not enough to simplify the graph, we needed to implement other simplification tactics.

First technic we used is *Node degree thresholding*. With this method we eliminated nodes with a low degree, considering they have less connections and therefore not crucial to the network's overall structure. The minimal number of edges a node must possess in order to be included in the condensed graph can be specified as a threshold (which is **5** in our case). After deleting nodes with less than five it is likely to have some nodes with zero degree, thus we eliminated them as well.

```

> reduced_email = igraph::delete.vertices(simple_email, igraph::degree(simple_email) < 5)
> reduced_email = igraph::delete.vertices(reduced_email, igraph::degree(reduced_email) == 0)
> vcount(reduced_email)
[1] 10155
> ecount(reduced_email)
[1] 118980

```

- `shortest.paths()`

The shortest pathways between each pair of vertices in a graph are determined using the `igraph::shortest.paths()` method. The shortest path between a source vertex and every

other vertex in the graph is determined using Dijkstra's method.


```
> reduced_email.sp = igraph::shortest.paths(reduced_email)
> reduced_email.sp
```

0	1	3	5	7	8	10	11	14	15	16	19	20	22	23	24	25	26	27	28	30	32	33	34	35	37	40	41	42	44	45
46	47	48	49	50	54	55	56	58	59	60	63	65	66	68	70	71	72	76	77	79	81	83	86	87	94	97	98	99	102	104
106	107	108	109	111	113	114	115	116	118	119	120	126	130	133	134	135	136	137	138	139	140	143	146	147	148	149	151	152	155	158
160	163	167	171	175	178	182	184	185	186	187	192	195	196	199	200	202	203	205	206	207	209	211	212	215	217	219	220	222	225	231
232	233	236	237	238	240	242	247	248	250	251	252	253	254	255	256	259	261	264	269	271	272	273	274	275	278	279	280	283	285	287
288	289	292	294	296	298	299	301	302	304	305	306	307	309	310	312	313	314	315	316	318	325	326	327	328	330	332	333	336	337	338
344	346	347	349	350	352	355	356	358	360	363	364	366	372	373	376	379	380	385	387	388	389	390	391	392	393	397	399	400	401	402
403	404	406	408	410	413	415	417	420	421	422	424	425	426	430	431	432	433	434	435	438	439	440	441	442	444	446	447	450	452	455
456	457	459	460	462	464	465	466	467	468	469	473	477	479	481	483	485	489	491	493	495	497	499	500	502	503	505	506	509	510	512
515	516	518	522	525	527	528	534	535	536	537	538	544	557	562	563	566	567	569	571	572	573	575	577	579	580	581	583	584	585	586
588	589	590	592	595	596	598	599	601	602	604	605	606	609	611	612	615	617	618	620	621	622	623	626	627	628	631	632	633	634	635
636	640	644	645	647	649	650	652	653	654	656	657	658	659	660	661	663	665	666	668	670	671	673	676	677	681	682	684	685	687	689

4. Explore other functions in the igraph package

- `mst()`


Minimal Spanning Tree connects all the vertices of the graph with the minimum possible total edge weight. Interpretation can vary depending on the context such as identifying most important (central nodes), efficient paths between pairs of vertices, and even identifying communities.



- `eigenvector_centrality()`

A measure of a vertex's relevance in a network called eigenvector centrality and considers both the vertex's connectivity to other nodes and the significance of the nodes to which those connections are made.

```
> ec <- igraph::evcent(reduced_email)
> plot(reduced_email, vertex.color=ec$vector, vertex.size=30*ec$vector, vertex.label=NA, layout=layout.circle)
>
```



- `summary()`


This function offers a summary of the graph, including its type, size, number of vertices and edges, and any properties.

```
> summary(reduced_email)
IGRAPH c51933d DN-- 10155 118980 --
+ attr: name (v/c)
> |
```

- Cluster_louvain (explained in more detail at 5.f)
- Layout_with_kk

This function from the igraph package uses the Kamada-Kawai technique to determine a graph's layout. By placing vertices in places that minimize the sum of the spring and electrical forces connecting them, the Kamada-Kawai algorithm, a force-directed graph layout technique, seeks to reduce the overall energy of the network.

```
> layout <- layout_with_kk(aggregated_graph, dim=2, kkconst=1)
> plot(aggregated_graph, layout=layout, vertex.label=NA, edge.arrow.size=0.5)
>
```



- Transitivity

This function determines a graph's transitivity, which is the proportion of triangles to

linked triples of vertices in the graph.

```
> transitivity(reduced_email)
[1] 0.06774965
```

- Triangles

This function determines a graph's triangle count, which is a gauge of the graph's clustering coefficient.

```
> triangles(reduced_email)
+ 740187/10155 vertices, named, from c51933d:
[1] 192 10 56 192 10 65 192 10 97 192 10 106 192 10 146 192 10 175
[19] 192 10 182 192 10 186 192 10 195 192 10 202 192 10 206 192 10 211
[37] 192 10 264 192 10 333 192 10 336 192 10 366 192 10 389 192 10 422
[55] 192 10 467 192 10 493 192 10 500 192 10 505 192 10 510 192 10 562
[73] 192 10 640 192 10 693 192 10 838 192 10 841 192 10 872 192 10 920
[91] 192 10 937 192 10 949 192 10 991 192 10 1010 192 10 1293 192 10 1370
[109] 192 10 1509 192 10 1518 192 10 1572 192 10 1623 192 10 1715 192 10 1844
[127] 192 10 1893 192 10 1955 192 10 3507 192 10 4272 192 10 4527 192 10 5357
[145] 192 10 5401 192 10 5718 192 10 5809 192 10 6904 192 10 7296 192 10 8344
[163] 192 10 9394 192 10 13972 192 10 19125 192 10 37994 192 10 40215 192 10 47471
```

- Isomorphic

This function determines if two graphs are isomorphic, which implies they share the same structure but may have distinct vertex and edge names. Output below is clearly false, because we eliminated some connections, resulting changing the structure.

```
> isomorphic(email_graph, reduced_email)
[1] FALSE
```

- Mean_distance

The term "mean distance" describes the shortest path's average length between every pair of vertices in a network. The smallest number of edges that must be crossed to get from one vertex to another is known as the shortest path length.

```
> mean_distance(reduced_email)
[1] 3.377349
```

- count_automorphisms()

This is a method to find the number of automorphisms in a graph, which are isomorphisms from the graph to itself. An automorphism of a graph is a permutation of its vertices that preserves its edges. Isomorphisms means, two graphs being structurally equivalent. It seems that there is a high degree of symmetry in the network.

```
> count_automorphisms(reduced_email)
$nof_nodes
[1] 1158

$nof_leaf_nodes
[1] 3

$nof_bad_nodes
[1] 0

$nof_canupdates
[1] 1

$max_level
[1] 358

$group_size
[1] "41193613532220680655717924052"
```

5. Determine the (a) central nodes(s) in the graph, (b) longest path(s), (c) largest clique(s), (d) ego(s), (e) power centrality, (f) find communities.

- central nodes(s) in the graph

There are several methods for calculating centrality in a network, and various centrality metrics can draw attention to various kinds of significant nodes.

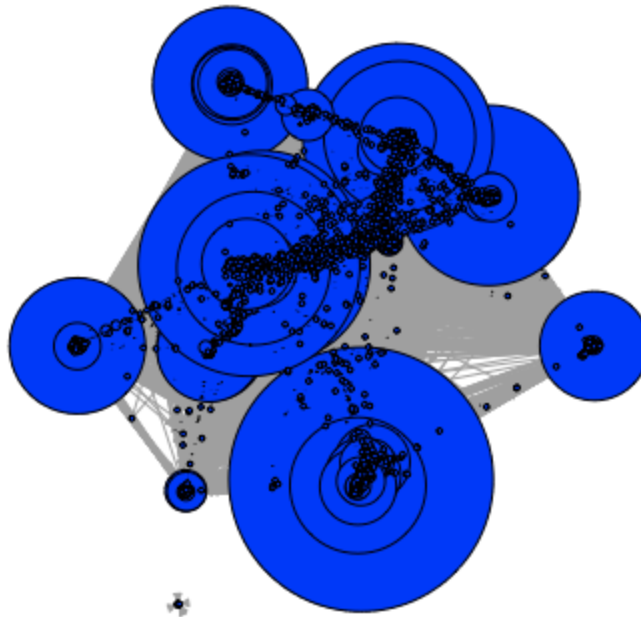
a. Betweenness centrality

This calculates how far a node is located along the network's shortest pathways to other nodes. Bridges connecting various regions of the network are built by nodes

with high betweenness centrality.

```
> reduced_email.betweenness <- centr_betw(reduced_email)
> node_size <- 100 * reduced_email.betweenness$res / max(reduced_email.betweenness$res)
> plot(reduced_email, vertex.size = node_size, vertex.color = "blue", vertex.label = NA, edge.arrow.size=0.3)
|
```

In order to visualize betweenness, we can filter nodes where size is proportional to the corresponding centrality score:




- b. Closeness centrality determined by taking the reciprocal of the lengths of the shortest routes that connect the node to every other node in the graph. As in betweenness centrality, we filter nodes by size as a measure of the importance of each node based on its closeness centrality, where larger nodes indicate greater importance.

```

> reduced_email_closeness <- centr_clo(reduced_email)
> if (any(!is.na(reduced_email_closeness$res))) {
+   node_size <- 20 * reduced_email_closeness$res
+   node_size[is.nan(node_size)] <- 0 # set NaN values to 0
+ } else {
+   node_size <- rep(0, vcount(reduced_email))
+ }
> plot(reduced_email, vertex.size = node_size, vertex.color = "red", vertex.label = NA, edge.arrow.size=0.3)

```



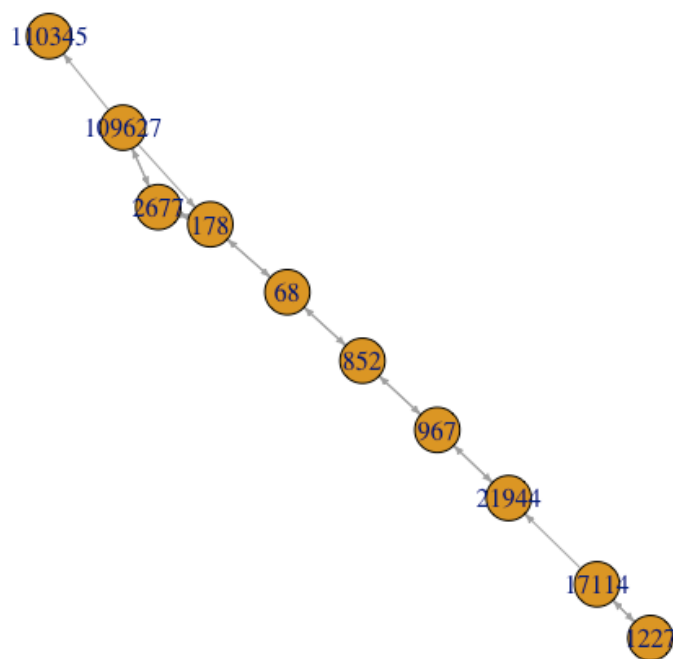
- longest path

In order to find the longest path of the graph, we can calculate the diameter of the graph, which is the length of the longest shortest path in the graph.

```

> diameter <- get_diameter(reduced_email)
> diameter
+ 10/10155 vertices, named, from c51933d:
[1] 1227 17114 21944 967 852 68 178 2677 109627 110345

```




- largest clique

This function returns all the maximal cliques in the graph, and we can select the largest one using the `which.max` function. After finding the largest clique, we can visualize it as subgraph.

```

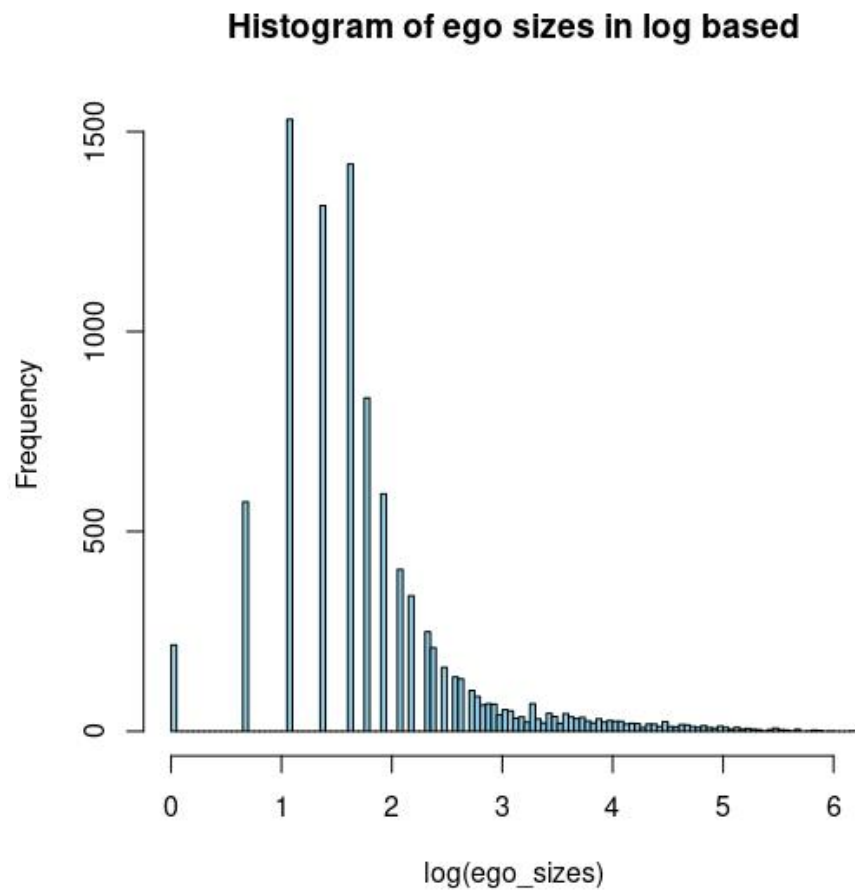
> largest_clique <- cliques[[which.max(lengths(cliques))]]
>
> largest_clique
+ 16/10155 vertices, named, from c51933d:
[1] 92476 146 336 147 406 158 505 118336 779 149 1515 1425 50 1 175 640

```




- (d) ego(s),

From the histogram below it is visible that it is right skewed which means generality of the nodes in the network have ego sizes and only among 10536 nodes 150 of them has ego sizes more than 100. The mean, median, maximum values are 11.2, 5, 479 respectively.



- (e) power centrality,

It is measure of the influence of nodes in the network using the connections of them. In this network we can clearly see from distribution that majority of the the values in the power centrality lies around zero which are very small number. This can indicate that few of the nodes in the network play crucial influence in the organization while lots of them do not have reasonable effect. The maximum value is 11.9 while median is 0.




- (f) find communities

We performed both clustering and community detection on the network. Community discovery is process of clustering nodes of the network which are densely connected but sparsely to the rest of the network. In order to understand the shared network characteristics among the nodes we used Clustering and Node aggregation technic. With this method, we find closely linked groups of nodes and consider them as a single node in the condensed graph. By doing so, it enables us to decrease the quantity of nodes and edges while maintaining the network's general structure.

We used Louvian method which is a well-known technique for community discovery in graphs. It seeks to divide a graph's nodes into distinct communities or clusters depending on the graph's topology. It is a two-stage procedure: the first phase involves assigning each node to a distinct community, and the second step involves iteratively merging neighboring communities to maximize a quality function that gauges the modularity of the final partition. The modularity gauges how closely related nodes within a community are to those connecting nodes between communities.

After aggregating the vertices which belong to the same community into a vertex, we removed the isolated vertices from the graph.




In the picture above we can clearly see that there are departments or teams in this network which have mostly close relationship with each other. One of the insight we can get from graph is that clusters “a”, “k”, “l”, “o” seems to be important having dense


relationship with almost other groups. Another striking pattern is that cluster “u” never emailed to any of the groups which may indicate that this group does not actively play role in the organization but only gets some updates from “b” and “m” teams. Next outstanding point can be that there are some groups such as “t”, “s”, “e”, “q” which are located at the edges of the network, seems to have close relationship with only specific departments.

6. Resulting graph with too many vertices and edges will look very messy in the plot. Try to filter vertices and their edges in some way having in resulting plot (visualization) 30 – 100 vertexes. Differentiate vertices (by color, size, shape) and edges (color, type) of graph. Think about opportunity to assign weights to edges differentiating them accordingly.

We used walktrap community detection on the network. It outputs the plot below:



We also cluster the reduced graph using louvain algorithm using resolution parameter 10 which produces 221 clusters.



We then assign weights to the edges and nodes to prune some of them. After selecting nodes with weight more than two using induced subgraph we end up with 34 nodes. We talked about interpretation in the last lection of the previous task so, we do not go to much detail here.

